

Homework 06 CSCI 036

Lucas Welch

Due: Friday, 2022-10-21

Instructions

Please box your answers. For numerical answers, this can be done using something like 34. For text answers, this can be done using something like My answer. The output of a code chunk is automatically boxed, so no need to do more.

Consider the following tibbles.

```
t1 <- tibble(
  emp_id = c(1001, 1116, 1239),
  class = c("temp", "full", "full")
)
t2 <- tibble(
  emp_id = c(1001, 1211, 1239, 1543),
  comp = c(16232, 42003, 51522, 44023)
)
```

- What is the primary key in `t1`, that is, a key that consists of a single variable?
- Write code to add the information from the `comp` variable of `t2` to the three observations in `t1`, but only with the rows of `t1`.
- Write code to remove the observations from `t1` where there are no `comp` values in `t2`. The columns of the output should be the same as those of `t1`.

a. `emp_id`

b.

```
t1 |>
  left_join(t2)
```

```
## Joining, by = "emp_id"
```

```
## # A tibble: 3 × 3
##   emp_id class  comp
##   <dbl> <chr> <dbl>
## 1   1001 temp  16232
## 2   1116 full    NA
## 3   1239 full  51522
```

c.

```
t1 |>
  semi_join(t2)
```

```
## Joining, by = "emp_id"
```

```
## # A tibble: 2 × 2
##   emp_id class
##   <dbl> <chr>
## 1   1001 temp
## 2   1239 full
```

#semi_join: It keeps only the values from `t1` that have values in `t2`

Consider the `nycflights13` package.

```
library(nycflights13)
```

The `airports` dataset contains the FAA designation of airports, along with their name and the latitude and longitude of the destination.

```
airports |> head()
```

```
## # A tibble: 6 × 8
##   faa   name          lat   lon   alt   tz dst  tzone
##   <chr> <chr>         <dbl> <dbl> <dbl> <dbl> <chr> <chr>
## 1 04G   Lansdowne Airport  41.1 -80.6  1044   -5 A   America/Ne...
## 2 06A   Moton Field Municipal Airport  32.5 -85.7   264   -6 A   America/Ch...
## 3 06C   Schaumburg Regional  42.0 -88.1   801   -6 A   America/Ch...
## 4 06N   Randall Airport    41.4 -74.4   523   -5 A   America/Ne...
## 5 09J   Jekyll Island Airport  31.1 -81.4    11   -5 A   America/Ne...
## 6 0A9   Elizabethton Municipal Airport  36.4 -82.2  1593   -5 A   America/Ne...
```

The `flights` dataset contains data for over 300,000 flights leaving the New York metropolitan area.

```
flights |> head()
```

```
## # A tibble: 6 × 19
##   year month   day dep_time sched_dep...1 dep_d...2 arr_t...3 sched...4 arr_d...5 carrier
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>      <int>      <dbl> <chr>
## 1  2013     1     1     517           515        2      830       819        11 UA
## 2  2013     1     1     533           529        4      850       830        20 UA
## 3  2013     1     1     542           540        2      923       850        33 AA
## 4  2013     1     1     544           545       -1     1004      1022       -18 B6
## 5  2013     1     1     554           600       -6      812       837       -25 DL
## 6  2013     1     1     554           558       -4      740       728        12 UA
## # ... with 9 more variables: flight <int>, tailnum <chr>, origin <chr>,
## #   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>, and abbreviated variable names 1sched_dep_time,
## #   2dep_delay, 3arr_time, 4sched_arr_time, 5arr_delay
```

Consider just keeping the origin and dest information along with a key for each flight.

```
origin_dest <-
  flights |>
  select(year:day, origin, dest)
```

Note that the name of the key in the `airport` dataset is `faa`, while the airport name might be either origin or dest. Still, a `left_join` can be used by setting the `by` parameter correctly. For instance, to join the longitude and latitude of the destination airport of the flight, use:

```
origin_dest |>
  left_join(airports, by = c("dest" = "faa")) |>
  head()
```

```
## # A tibble: 6 × 12
##   year month   day origin dest  name          lat   lon   alt   tz dst  tzone
##   <int> <int> <int> <chr> <chr> <chr>        <dbl> <dbl> <dbl> <dbl> <chr> <chr>
## 1  2013     1     1 EWR   IAH   George Bus... 30.0 -95.3   97   -6 A    Amer...
## 2  2013     1     1 LGA   IAH   George Bus... 30.0 -95.3   97   -6 A    Amer...
## 3  2013     1     1 JFK   MIA   Miami Intl    25.8 -80.3    8   -5 A    Amer...
## 4  2013     1     1 JFK   BQN   <NA>          NA    NA    NA    NA <NA> <NA>
## 5  2013     1     1 LGA   ATL   Hartsfield... 33.6 -84.4  1026  -5 A    Amer...
## 6  2013     1     1 EWR   ORD   Chicago Oh... 42.0 -87.9   668  -6 A    Amer...
```

Modify the above code to join the longitude and latitude of the origin airport for flights in `origin_dest`.

```
origin_dest |>
  left_join(airports, by = c("origin" = "faa")) |>
  head()
```

```
## # A tibble: 6 × 12
##   year month   day origin dest  name          lat   lon   alt   tz dst  tzone
##   <int> <int> <int> <chr> <chr> <chr>        <dbl> <dbl> <dbl> <dbl> <chr> <chr>
## 1  2013     1     1 EWR   IAH   Newark Lib... 40.7 -74.2   18   -5 A    Amer...
## 2  2013     1     1 LGA   IAH   La Guardia   40.8 -73.9   22   -5 A    Amer...
## 3  2013     1     1 JFK   MIA   John F Ken... 40.6 -73.8   13   -5 A    Amer...
## 4  2013     1     1 JFK   BQN   John F Ken... 40.6 -73.8   13   -5 A    Amer...
## 5  2013     1     1 LGA   ATL   La Guardia   40.8 -73.9   22   -5 A    Amer...
## 6  2013     1     1 EWR   ORD   Newark Lib... 40.7 -74.2   18   -5 A    Amer...
```

(Exercise 13.6.1.1 of R for Data Science by Golemund & Wickham)

- a. Write code to compute the average flight arrival delay from the `flights` dataset (removing NA values) grouped by destination airport.
- b. The following code draws a map of the United States with the airports that are destinations in the `flights` dataset.

```
flights |>
  select(dest) |>
  unique() |>
  left_join(select(airports, faa, lat, lon),
            by = c("dest" = "faa")) |>
  filter(lon > -140) |>
  ggplot(aes(lon, lat)) +
    borders("state") +
    geom_point() +
    theme_void() +
    coord_quickmap()
```



Note that by using `unique()`, each airport only appears once in the tibble being plotted, which makes the number of rows 105 instead of over 330,000. This speeds up the graphic plotting by a lot! Modify this code to color the airports in this plot based on their average flight delay.

a.

```
flights |>
  group_by(dest) |>
  mutate(on_time_flights = arr_delay >= 0 ) |>
  summarize(avg_on_time = mean(on_time_flights, na.rm = TRUE))
```


```
## # A tibble: 105 × 2
##   dest   avg_on_time
##   <chr>     <dbl>
## 1 ABQ       0.429
## 2 ACK       0.413
## 3 ALB       0.452
## 4 ANC       0.625
## 5 ATL       0.493
## 6 AUS       0.419
## 7 AVL       0.487
## 8 BDL       0.359
## 9 BGR       0.397
## 10 BHM      0.465
## # ... with 95 more rows
```

b.

```
flights |>
  select(dest, arr_delay) |>
  unique() |>
  left_join(select(airports, faa, lat, lon),
            by = c("dest" = "faa")) |>
  group_by(dest) |>
  mutate(on_time_flights = arr_delay >= 0 ) |>
  select(arr_delay, on_time_flights, lon, lat) |>
  mutate(avg_on_time = mean(on_time_flights, na.rm = TRUE)) |>

  filter(lon > -140) |>
  ggplot(aes(lon, lat, color = avg_on_time)) +
    borders("state") +
    geom_point() +
    theme_void() +
    coord_quickmap()
```

```
## Adding missing grouping variables: `dest`
```



#COME BACK REVIEW

(Exercise 13.6.1.1 of R for Data Science by Golemund & Wickham)

Write code to add an artificial (surrogate) key to `flights` with the `row_number` function which is the first column of the new tibble.

```
flights |>
  mutate(key = row_number(), .before = 1)
```

```
## # A tibble: 336,776 × 20
##       key  year month   day dep_time sched_dep_...1 dep_d...2 arr_t...3 sched...4 arr_d...5
##   <int> <int> <int> <int>   <int>         <int>      <dbl>   <int>    <int>    <dbl>
## 1     1     2013     1     1     517           515        2     830     819     11
## 2     2     2013     1     1     533           529        4     850     830     20
## 3     3     2013     1     1     542           540        2     923     850     33
## 4     4     2013     1     1     544           545       -1    1004    1022    -18
## 5     5     2013     1     1     554           600       -6     812     837    -25
## 6     6     2013     1     1     554           558       -4     740     728     12
## 7     7     2013     1     1     555           600       -5     913     854     19
## 8     8     2013     1     1     557           600       -3     709     723    -14
## 9     9     2013     1     1     557           600       -3     838     846     -8
## 10    10    2013     1     1     558           600       -2     753     745      8
## # ... with 336,766 more rows, 10 more variables: carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>, and abbreviated variable names
## #   1sched_dep_time, 2dep_delay, 3arr_time, 4sched_arr_time, 5arr_delay
```


(Exercise 13.6.1.1 of R for Data Science by Golemund & Wickham modified.)

- a. Create a tibble that counts the number of times that any particular plane has flown—removing those observations where the tail number is unknown. Filter this tibble to only include planes that have flown at least 100 times.
- b. Using a semi-join, filter `flights` to only show flights with planes that have flown at least 100 flights.

a.

```
TB <- flights |>
  filter(!is.na(tailnum)) |>
  group_by(tailnum) |>
  summarize(count = n()) |>
  filter(count >= 100)
```

TB

```
## # A tibble: 1,217 × 2
##   tailnum count
##   <chr>   <int>
## 1 N0EGMQ    371
## 2 N10156    153
## 3 N10575    289
## 4 N11106    129
## 5 N11107    148
## 6 N11109    148
## 7 N11113    138
## 8 N11119    148
## 9 N11121    154
## 10 N11127   124
## # ... with 1,207 more rows
```

b.

```
flights |>
  semi_join(TB)
```

```
## Joining, by = "tailnum"
```

```
## # A tibble: 228,390 × 19
##   year month   day dep_time sched_de...1 dep_d...2 arr_t...3 sched...4 arr_d...5 carrier
##   <int> <int> <int>   <int>      <int>      <dbl>    <int>    <int>    <dbl> <chr>
## 1  2013     1     1     517        515         2      830      819      11 UA
## 2  2013     1     1     533        529         4      850      830      20 UA
## 3  2013     1     1     544        545        -1     1004     1022     -18 B6
## 4  2013     1     1     554        558        -4      740      728      12 UA
## 5  2013     1     1     555        600        -5      913      854      19 B6
## 6  2013     1     1     557        600        -3      709      723     -14 EV
## 7  2013     1     1     557        600        -3      838      846      -8 B6
## 8  2013     1     1     558        600        -2      849      851      -2 B6
## 9  2013     1     1     558        600        -2      853      856      -3 B6
## 10 2013     1     1     558        600        -2      923      937     -14 UA
## # ... with 228,380 more rows, 9 more variables: flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>, and abbreviated variable names
## #   1sched_dep_time, 2dep_delay, 3arr_time, 4sched_arr_time, 5arr_delay
```

State which of the following

- `pivot_wider`
- `pivot_longer`
- `separate`
- `unite`

is the right function to use in the situation.

- Two variable values need to be combined into one variable entry.
- The entries of a column should be column names instead.
- An entry contains values for two or more variables.
- The column names are measurements rather than something you measure.

a.

b.

c.

d.

Consider the following dataset that gives the number of days of rainfall for five cities over three months.

```
df1 <- read_csv2('City;January;February;March
Atlanta, Georgia;11;10;10
Austin, Texas;7;7;9
Baltimore, Maryland;10;9;10
Birmingham, Alabama;11;10;10
Boston, Massachusetts;11;10;12')
df1
```

```
## # A tibble: 5 × 4
##   City          January February March
##   <chr>          <dbl>    <dbl> <dbl>
## 1 Atlanta, Georgia      11        10    10
## 2 Austin, Texas         7         7     9
## 3 Baltimore, Maryland  10         9    10
## 4 Birmingham, Alabama  11        10    10
## 5 Boston, Massachusetts 11        10    12
```

a. Tidy this data.

b. Find the mean number of days of rainfall from January through March for each of the five cities.

a.

```
df1 |>
  separate(City, into = c("city", "states"))
```

```
## # A tibble: 5 × 5
##   city      states      January February March
##   <chr>    <chr>      <dbl>    <dbl> <dbl>
## 1 Atlanta Georgia      11        10    10
## 2 Austin  Texas         7         7     9
## 3 Baltimore Maryland    10         9    10
## 4 Birmingham Alabama     11        10    10
## 5 Boston  Massachusetts 11        10    12
```

b.

```
df1 |>
  pivot_longer('January':'February':'March', names_to = "month", values_to = "rain_value
s") |>
  group_by(City) |>
  summarize(total_rain_fall = mean(rain_values))
```

```
## Warning in x:y: numerical expression has 2 elements: only the first used
```

```
## # A tibble: 5 × 2
##   City                total_rain_fall
##   <chr>                <dbl>
## 1 Atlanta, Georgia    10.3
## 2 Austin, Texas       7.67
## 3 Baltimore, Maryland 9.67
## 4 Birmingham, Alabama 10.3
## 5 Boston, Massachusetts 11
```

Now suppose the data from the last problem was presented in the following CSV file:

```
df2 <- read_csv2("
Atlanta, Georgia;Jan/11
Atlanta, Georgia;Feb/10
Atlanta, Georgia;Mar/10
Austin, Texas;Jan/7
Austin, Texas;Feb/7
Austin, Texas;Mar/9
Baltimore, Maryland;Jan/10
Baltimore, Maryland;Feb/9
Baltimore, Maryland;Mar/10
Birmingham, Alabama;Jan/11
Birmingham, Alabama;Feb/10
Birmingham, Alabama;Mar/10
Boston, Massachusetts;Jan/11
Boston, Massachusetts;Feb/10
Boston, Massachusetts;Mar/12",
col_names = c("City", "Days_rain")
)
```

a. Tidy this data.

b. Create a horizontal bar plot that for each of the three months, making five bars for the five cities where the height of each bar is the total number of days of rain for that city from January through March.

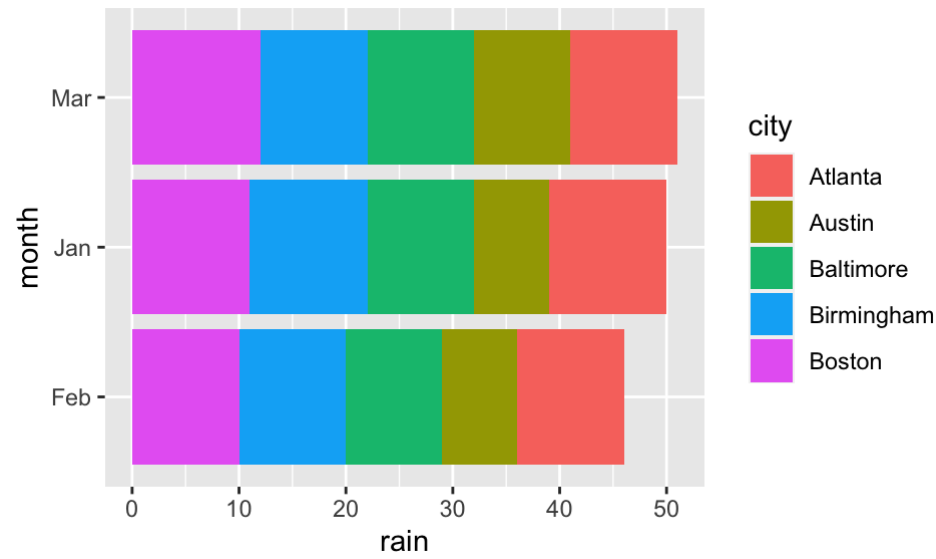
a.

```
df2 |>
  separate(City, into = c("city", "state")) |>
  separate(Days_rain, into = c("month", "rain")) |>
  pivot_wider(names_from = month, values_from = rain)
```

```
## # A tibble: 5 × 5
##   city      state      Jan   Feb   Mar
##   <chr>    <chr>    <chr> <chr> <chr>
## 1 Atlanta  Georgia    11    10    10
## 2 Austin   Texas      7      7      9
## 3 Baltimore Maryland  10      9    10
## 4 Birmingham Alabama    11     10    10
## 5 Boston   Massachusetts 11     10    12
```

b.

```
df2 |>
  separate(City, into = c("city", "state")) |>
  separate(Days_rain, into = c("month", "rain"), convert = TRUE) |>
  ggplot(aes(x = month, y = rain)) +
  geom_bar(aes(fill = city), stat = 'identity') +
  coord_flip()
```



Consider the following data set.

```
state_data <- tribble(
  ~state, ~type, ~number,
  "Oregon", "year_founded", 1859,
  "Oregon", "population_2020_census", 4237256,
  "Oregon", "population_2010_census", 3831074,
  "California", "year_founded", 1850,
  "California", "population_2020_census", 39538223,
  "Washington", "population_2020_census", 7705281,
  "Washington", "population_2010_census", 6724540,
)
```

- Are all the values under `number` measuring the same thing?
- Tidy this data.
- How many NA values are in your table?

a.

b.

```
state_data |>
  pivot_wider(names_from = type, values_from = number)
```

```
## # A tibble: 3 × 4
##   state      year_founded population_2020_census population_2010_census
##   <chr>          <dbl>              <dbl>              <dbl>
## 1 Oregon          1859                4237256            3831074
## 2 California      1850                39538223             NA
## 3 Washington      NA                  7705281            6724540
```

c.

Consider the following dataset. Bring together the first and last names into a single field `name`, where the last name goes first, followed by a comma, then a space, and finally the first name. For instance if the first name was Pete and the last name Aguilar, the result should be `Aguilar, Pete`.

```
jan6_committee <-
  tribble(
    ~first_name, ~last_name, ~membership,
    "Bennie", "Thompson", "Chairperson",
    "Zoe", "Lofgren", "Majority",
    "Adam", "Schiff", "Majority",
    "Pete", "Aguilar", "Majority",
    "Stephanie", "Murphy", "Majority",
    "Jamie", "Raskin", "Majority",
    "Elaine", "Luria", "Majority",
    "Liz", "Cheney", "Minority",
    "Adam", "Kinzinger", "Minority"
  )
jan6_committee
```

```
## # A tibble: 9 × 3
##   first_name last_name membership
##   <chr>      <chr>      <chr>
## 1 Bennie     Thompson  Chairperson
## 2 Zoe        Lofgren   Majority
## 3 Adam       Schiff    Majority
## 4 Pete       Aguilar   Majority
## 5 Stephanie  Murphy    Majority
## 6 Jamie      Raskin    Majority
## 7 Elaine     Luria     Majority
## 8 Liz        Cheney    Minority
## 9 Adam       Kinzinger Minority
```

```
jan6_committee |>
  unite(col = full_name, last_name, first_name, sep = ",")
```

```
## # A tibble: 9 × 2
##   full_name      membership
##   <chr>          <chr>
## 1 Thompson,Bennie Chairperson
## 2 Lofgren,Zoe     Majority
## 3 Schiff,Adam     Majority
## 4 Aguilar,Pete    Majority
## 5 Murphy,Stephanie Majority
## 6 Raskin,Jamie     Majority
## 7 Luria,Elaine    Majority
## 8 Cheney,Liz      Minority
## 9 Kinzinger,Adam  Minority
```